

Introduction to Reinforcement Learning

github.com/kengz/openai_lab

*Wah Loon Keng
Laura Graesser*

9 Dec 2017

RL system

- Agent (OpenAI Lab)
- Environment (OpenAI gym)
- SARS (state, action, reward, state_next) time transition

Characteristics

- Goal oriented, reward signals
- On-line, interactive
- Stateful environment (MDP)
- No full access to the function to optimize (no labeled data)




Image from http://wiki.ubc.ca/Course:CPSC522/Reinforcement_Learning

RL as MDP

Can be specified with:

S: state space (discrete/cont.)

A: action space (discrete/cont.)

R: reward distribution, $(\cdot|s,a)$

P: transition probability, $(\cdot|s,a)$

MDP:

1. Init, $t = 0$, env samples state $s_0 \sim p(s_0)$
2. For $t = 0$ until done:
 - 2.1. Agent selects action, a_t
 - 2.2. Env samples reward, $r_t \sim R(\cdot|s_t, a_t)$
 - 2.3. Env samples next state, $s_{t+1} \sim P(\cdot|s_t, a_t)$
 - 2.4. Agent receives r_t, s_{t+1}

CartPole-v0

State space (continuous, 4 dims):

- position
- velocity
- angle
- angular velocity
- e.g.: [0.023, 0.004, -0.061, 0.028]

Action space (discrete, 2 actions):

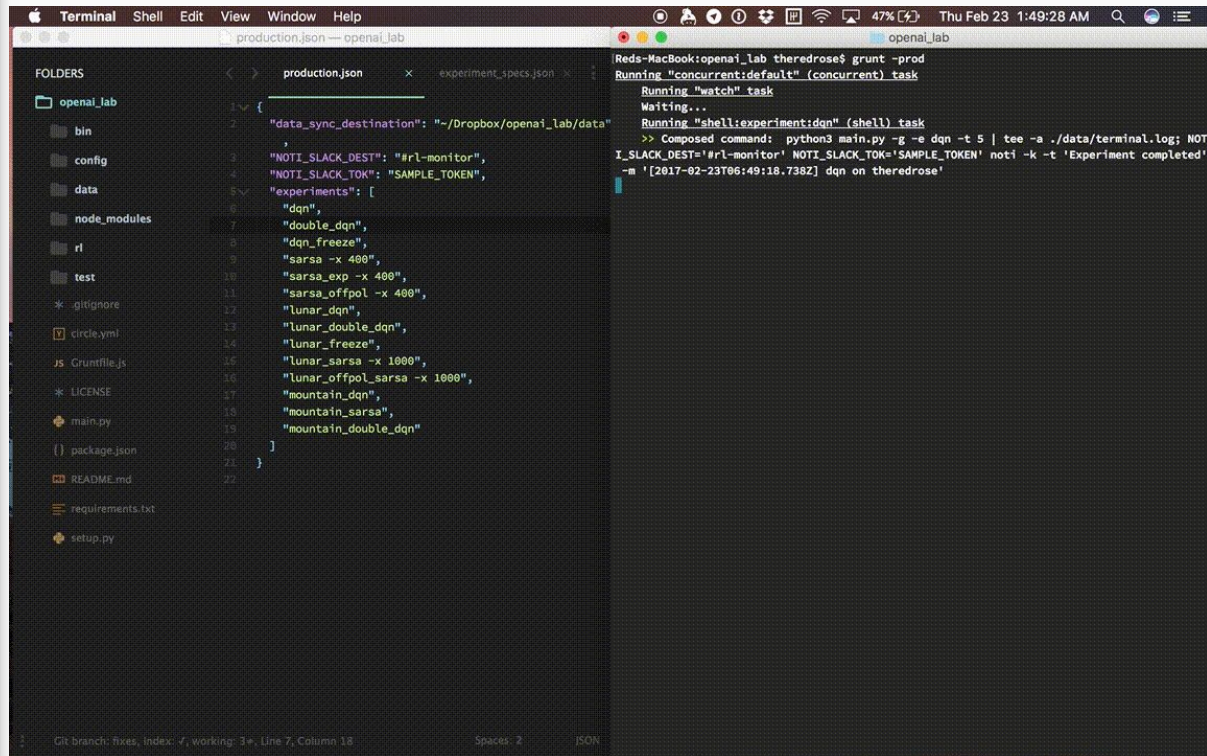
- push left: 0
- push right: 1
- e.g.: 0

Reward: +1 per timestep for staying up

Solution/epi: total rewards = 200

Solution: mean rewards > 195.0 over 100 consecutive episodes

github.com/kengz/openai_lab



The screenshot shows a terminal window with a file explorer on the left and a terminal on the right. The file explorer shows the contents of the 'openai_lab' directory, including 'bin', 'config', 'data', 'node_modules', 'rl', 'test', '.gitignore', 'circle.yml', 'Gruntfile.js', 'LICENSE', 'main.py', 'package.json', 'README.md', 'requirements.txt', and 'setup.py'. The terminal window shows the command 'production.json' being executed, which outputs a JSON configuration for the 'production' task. The configuration includes 'data_sync_destination', 'NOTI_SLACK_DEST', 'NOTI_SLACK_TOKEN', and a list of experiments: 'dqn', 'double_dqn', 'dqn_freeze', 'sarsa -x 400', 'sarsa_exp -x 400', 'sarsa_offpol -x 400', 'lunar_dqn', 'lunar_double_dqn', 'lunar_freeze', 'lunar_sarsa -x 1000', 'lunar_offpol_sarsa -x 1000', 'mountain_dqn', 'mountain_sarsa', and 'mountain_double_dqn'. The terminal also shows the command 'experiment_specs.json' being executed, which outputs a JSON configuration for the 'experiment_specs' task. The configuration includes 'data_sync_destination', 'NOTI_SLACK_DEST', 'NOTI_SLACK_TOKEN', and a list of experiments: 'dqn', 'double_dqn', 'dqn_freeze', 'sarsa -x 400', 'sarsa_exp -x 400', 'sarsa_offpol -x 400', 'lunar_dqn', 'lunar_double_dqn', 'lunar_freeze', 'lunar_sarsa -x 1000', 'lunar_offpol_sarsa -x 1000', 'mountain_dqn', 'mountain_sarsa', and 'mountain_double_dqn'. The terminal also shows the command 'grunt -prod' being executed, which outputs 'Running "concurrent:default" (concurrent) task' and 'Running "watch" task'.

```
production.json
{
  "data_sync_destination": "~/Dropbox/openai_lab/data",
  "NOTI_SLACK_DEST": "#rl-monitor",
  "NOTI_SLACK_TOKEN": "SAMPLE_TOKEN",
  "experiments": [
    "dqn",
    "double_dqn",
    "dqn_freeze",
    "sarsa -x 400",
    "sarsa_exp -x 400",
    "sarsa_offpol -x 400",
    "lunar_dqn",
    "lunar_double_dqn",
    "lunar_freeze",
    "lunar_sarsa -x 1000",
    "lunar_offpol_sarsa -x 1000",
    "mountain_dqn",
    "mountain_sarsa",
    "mountain_double_dqn"
  ]
}

experiment_specs.json
{
  "data_sync_destination": "~/Dropbox/openai_lab/data",
  "NOTI_SLACK_DEST": "#rl-monitor",
  "NOTI_SLACK_TOKEN": "SAMPLE_TOKEN",
  "experiments": [
    "dqn",
    "double_dqn",
    "dqn_freeze",
    "sarsa -x 400",
    "sarsa_exp -x 400",
    "sarsa_offpol -x 400",
    "lunar_dqn",
    "lunar_double_dqn",
    "lunar_freeze",
    "lunar_sarsa -x 1000",
    "lunar_offpol_sarsa -x 1000",
    "mountain_dqn",
    "mountain_sarsa",
    "mountain_double_dqn"
  ]
}

grunt -prod
Running "concurrent:default" (concurrent) task
Running "watch" task
Waiting...
Running "shell:experiment:dqn" (shell) task
>> Composed command: python3 main.py -g -e dqn -t 5 | tee -a ./data/terminal.log; NOTI_SLACK_DEST='#rl-monitor' NOTI_SLACK_TOKEN='SAMPLE_TOKEN' noti -k -t 'Experiment completed' -m '[2017-02-23T06:49:18.738Z] dqn on theredrose'
```

What to learn in an RL problem?

1. What action to take? → **Policy** based algorithms
2. How good actions are given states? → **Value** based algorithms
3. The dynamics of the system? → **Model** based algorithms

Some or all of the above?

Deep-RL hint: use deep neural networks to approximate these functions.

A policy is a function $\pi : S \rightarrow A$

Find a policy π^* that max. the cum. discounted reward $\sum_{t \geq 0} \gamma^t r_t$

$$\pi^* = \arg \max_{\pi} E \left[\sum_{t \geq 0} \gamma^t r_t | \pi \right]$$

with $s_0 \sim p(s_0)$, $a_t \sim \pi(\cdot | s_t)$, $s_{t+1} \sim p(s_t, a_t)$

Value function: how good is a state?

$$V^{\pi}(s) = E \left[\sum_{t \geq 0} \gamma^t r_t | s_0 = s, \pi \right]$$

Q-value function: how good is a state-action pair?

$$Q^{\pi}(s, a) = E \left[\sum_{t \geq 0} \gamma^t r_t | s_0 = s, a_0 = a, \pi \right]$$

Q-learning

Use a neural net to approximate the Q-function

$$Q^\pi(s, a) = E\left[\sum_{t \geq 0} \gamma^t r_t \mid s_0 = s, a_0 = a, \pi\right]$$

Neural net approx. with param θ : $Q(s, a; \theta) \approx Q^*(s, a)$

With a perfect Q function we know how to act optimally - policy for free

$$\pi^*(s) = \max_a Q(s, a)$$

How to learn it?
Bootstrapping

Bellman equation for optimal policy:

$$Q^*(s, a) = r(s, a) + \gamma E[V(s' \mid s, a)]$$

$$Q^*(s, a) \approx r(s, a) + \gamma \max_{a'} Q^*(s', a') \mid s, a$$

Algorithm DQN

For $i = 1 \dots N$:

Gather data (s_i, a_i, r_i, s'_i) by acting in the environment using some policy
for $j = 1 \dots K$:

1. Calculate target values for each example

$$y_i = r_i + \gamma \max_{a'} Q(s'_i, a'; \theta_{i-1}) | s_i, a_i$$

2. Update network parameters, using MSE loss

$$L_j(\theta) = \frac{1}{2} \sum_i ||(y_i - Q(s_i, a_i; \theta_i))||^2$$

Properties

- Value-based
- Off-policy
- Model-free
- Online/Offline

Pros & Cons

Pros:

- Low variance - Q function estimates the expected cumulative discounted rewards from (s^t, a^t)
- Fairly sample efficient

Cons:

- High bias (critic imperfect)
- Q function has no convergence guarantee
- Exploration problem

Improving Q-learning

- Replay memory: somewhat decorrelates data, more sample efficient
- Target networks: make the moving target move less
- Different networks to select next state action and to evaluate it

Policy Gradient

Policy Gradient Formalism

Define a class of parametrized policies, $\Pi = \{\pi_\theta, \theta \in \mathbb{R}^m\}$

For each π_θ , define its value $J(\theta) = E\left[\sum_{t \geq 0} \gamma^t r_t | \pi_\theta\right]$

Find the optimal policy $\theta^* = \arg \max_{\theta} J(\theta)$

For trajectory $\tau = (s_0, a_0, r_0, s_1, \dots)$ and scalar score function $f(\tau)$

$$J(\theta) = E_{\tau \sim p(\tau; \theta)} [f(\tau)] = \int_{\tau} f(\tau) p(\tau; \theta) d\tau$$

Weights the probability for action a given s

Gradient update for convergence (with some magic):

$$\begin{aligned} \nabla_{\theta} J(\theta) &= E_{\tau \sim p(\tau; \theta)} [f(\tau) \nabla_{\theta} \log p(\tau; \theta)] \\ &\approx \sum_{t \geq 0} f(\tau) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \end{aligned}$$

Policy Gradient Algorithm (REINFORCE, Williams 1992)

Algorithm REINFORCE:

Initialize weights θ , learning rate α

for each episode (trajectory) $\tau = \{s_0, a_0, r_0, s_1, \dots, r_T\} \sim \pi_\theta$

for $t = 0$ to T do

$$\theta \leftarrow \theta + \alpha f(\tau)_t \nabla_\theta \log \pi_\theta(a_t | s_t)$$

end for

end for

Properties

- Policy-based
- On-policy
- Model-free
- Offline (samples whole episodes)

Pros & Cons

Pros:

- Unbiased, uses direct sampling

Cons:

- High variance since the sampled reward varies per episode, credit assignment
- Inefficient sampling

Improving REINFORCE by lowering the variance:

Given $\nabla_{\theta} J(\theta) \approx \sum_{t \geq 0} f(\tau) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$, improve baseline with:

1. reward as weightage $f(\tau) = \sum_{t' \geq t} r_{t'}$

2. add discount factor $f(\tau) = \sum_{t' \geq t} \gamma^{t'-t} r_{t'}$

3. introduce baseline $f(\tau) = \sum_{t' \geq t} \gamma^{t'-t} r_{t'} - b(s_t)$

4. advantage function $f(\tau) = Q^{\pi}(s_t, a_t) - V^{\pi}(s_t) = A^{\pi}(s_t, a_t)$

$$\nabla_{\theta} J(\theta) \approx \sum_{t \geq 0} (Q^{\pi_{\theta}}(s_t, a_t) - V^{\pi_{\theta}}(s_t)) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$

Actor-Critic

What to fit for the critic?

A, Q, or V?

$$\nabla_{\theta} J(\theta) \approx \sum_{t \geq 0} (Q^{\pi_{\theta}}(s_t, a_t) - V^{\pi_{\theta}}(s_t)) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$

$$A^{\pi}(s_t, a_t) = Q^{\pi}(s_t, a_t) - V^{\pi}(s_t)$$

Q can be approximated with the reward from the current state and action, and V

$$Q^{\pi}(s_t, a_t) = r(s_t, a_t) + E_{s_{t+1} \sim p(s_{t+1} | s_t, a_t)} [V^{\pi}(s_{t+1})]$$

$$Q^{\pi}(s_t, a_t) \approx r(s_t, a_t) + V^{\pi}(s_{t+1})$$

$$A^{\pi}(s_t, a_t) = r(s_t, a_t) + V^{\pi}(s'_t) - V^{\pi}(s_t)$$

We can just learn V

How to learn V? The bootstrap again!

$$V(s_t) = r(s_t, a_t) + V(s'_t)$$

Actor-Critic Algorithm

For $i = 1 \dots N$:

1. Gather data (s_i, a_i, r_i, s'_i) by acting in the environment using your policy

2. Update V

for $j = 1 \dots K$:

Calculate target values for each example

$$y_i = r_i + V(s'_i)$$

Update network parameters, using MSE loss

$$L_j(\theta) = \frac{1}{2} \sum_i \|(y_i - V(s_i; \theta_i))\|^2$$

3. Evaluate A , $A^\pi(s_i, a_i) = r(s_i, a_i) + V^\pi(s'_i) - V^\pi(s_i)$

4. Calculate gradient, $\nabla_\phi J(\phi) \approx \sum_i A_t^\pi(s_i, a_i) \nabla_\phi \log \pi_\phi(a_i | s_i)$

5. Use gradient to update parameters ϕ

Properties

- Value and policy -based
- On-policy
- Model-free
- Online/offline

Pros & Cons

Pros:

- Low bias
- Low variance

Cons:

- Sample inefficient (from the policy gradient)

OpenAI Lab Demo

Appendix

Experience Replay

Memory unit

At each time step store $\langle S, A, R, S' \rangle$

- State, S
- Action taken in S
- Reward after taking action A
- Next state, S'

Parameters

- Max memory length

Advantages

- Can learn from past experiences as well as present experience
- Can learn about optimal policy whilst following an exploratory policy
- Decorrelates data (somewhat)